

```
## Comcast Telecom Complaints Dataset
```

## Context

Estimation of Comcast Customer Top Complaints.

## Acknowledgements

Kaggle Datasets

```
import warnings
warnings.filterwarnings(action='ignore')

import plotly.express as px
from plotly.subplots import make_subplots
import plotly.graph_objects as go

import seaborn as sns
import matplotlib.pyplot as plt
from matplotlib import rcParams
```

```
####! python -m pip install spacy==3.0.0
```

```
####! pip uninstall numpy -y
```

```
####! pip install numpy==1.18.5
```

```
# Install libraries
#!pip install boostaroota==1.3
#!pip install h2o==3.36.1.3
#!pip install ppscore==3.9.0
#!pip install imblearn
```

```
#!pip install optuna
```

```
####! pip install shap
####! pip install catboost
```

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import os
from ipywidgets import interact, interactive, fixed, interact_manual
import ipywidgets as widgets
import plotly.express as px
import matplotlib.pyplot as plt
import plotly.graph_objs as go
from tqdm import tqdm
from sklearn.metrics import mean_squared_error
import tensorflow as tf
from sklearn import model_selection as sk_model_selection
from xgboost.sklearn import XGBRegressor
from sklearn.metrics import mean_squared_error, roc_auc_score, precision_score
from sklearn import metrics
import optuna
from boostaroota import BoostARoota
from sklearn.metrics import log_loss
from optuna.samplers import TPESampler
import functools
from functools import partial
```

```
import xgboost as xgb
import joblib
from matplotlib_venn import venn2, venn2_circles, venn2_unweighted
from matplotlib_venn import venn3, venn3_circles
import statsmodels.api as sm
import pylab
from xgboost import plot_tree
import shap
from xgboost.sklearn import XGBClassifier
from sklearn.metrics import mean_squared_error, roc_auc_score, precision_score
from sklearn import metrics
from sklearn.metrics import log_loss
from sklearn.metrics import confusion_matrix, recall_score, precision_score, precision_recall_curve, auc, f1_score, \
    average_precision_score, accuracy_score, roc_curve
from sklearn.preprocessing import LabelEncoder
import h2o
from h2o.automl import H2OAutoML
from catboost import Pool, CatBoostRegressor, cv
import tensorflow as tf
from tensorflow.keras.utils import plot_model
from tensorflow.keras.models import Model, load_model
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.callbacks import ModelCheckpoint
from tensorflow.keras.callbacks import ReduceLROnPlateau
from tensorflow.keras.layers import BatchNormalization
from tensorflow.keras.layers import Dense, Dropout, Input
from tensorflow.keras.layers import Concatenate, LSTM, GRU
from tensorflow.keras.layers import Bidirectional, Multiply
import seaborn as sns
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
```

```
SEED = 42
```

```
import numpy as np
import pandas as pd
```

```
####!mkdir ~/.kaggle
```

```
####!cp /kaggle.json ~/.kaggle/
```

```
####!chmod 600 ~/.kaggle/kaggle.json
```

```
####!pip install keras-tuner
```

```
####!pip install kaggle
```

```
####! kaggle datasets download -d yasserh/comcast-telecom-complaints
```

```
####! unzip /content/comcast-telecom-complaints.zip
```

```
comcast = pd.read_csv("/content/Comcast.csv")
```

```
training_data = comcast.sample(frac=0.7, random_state=25)
testing_data = comcast.drop(training_data.index)
```

```
print(training_data.shape, testing_data.shape)
```

```
(1557, 11) (667, 11)
```

```
import nltk
nltk.download('punkt')
```

```
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data]   Package punkt is already up-to-date!
True
```

```
####! python -m textblob.download_corpora
```

```
####! pip install tensorflow==2.8.0
```

```
import tensorflow as tf
```

```
import pandas, numpy, string, textblob
import pickle
from sklearn import model_selection, preprocessing, linear_model, naive_bayes, metrics, svm, decomposition, ensemble
from sklearn.feature_extraction.text import TfidfVectorizer, CountVectorizer
import xgboost
from keras import layers, models, optimizers
from keras.preprocessing import text, sequence
import matplotlib.pyplot as plt
```

```
import pandas as pd
from xgboost import XGBRegressor as xgb
from sklearn.model_selection import GridSearchCV
import numpy as np
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings("ignore")
from sklearn.metrics import mean_squared_error
import sklearn
import seaborn as sns
from sklearn.model_selection import KFold
```

```
print(training_data.shape, testing_data.shape)
```

```
(1557, 11) (667, 11)
```

```
training_data.columns
```

```
Index(['Ticket #', 'Customer Complaint', 'Date', 'Date_month_year', 'Time',
      'Received Via', 'City', 'State', 'Zip code', 'Status',
      'Filing on Behalf of Someone'],
      dtype='object')
```

```
testing_data.columns
```

```
Index(['Ticket #', 'Customer Complaint', 'Date', 'Date_month_year', 'Time',
      'Received Via', 'City', 'State', 'Zip code', 'Status',
      'Filing on Behalf of Someone'],
      dtype='object')
```

```
training_data.to_csv("train.csv")
```

```
testing_data.to_csv("test.csv")
```

```
####!pip install -q -U watermark
```

```
####!pip install -qq transformers
```

## ▼ Setup & Config

```
##@title Setup & Config
import transformers
from transformers import BertModel, BertTokenizer, AdamW, get_linear_schedule_with_warmup
import torch

import numpy as np
import pandas as pd
import seaborn as sns
from pylab import rcParams
import matplotlib.pyplot as plt
from matplotlib import rc
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report
from collections import defaultdict
from textwrap import wrap

from torch import nn, optim
from torch.utils.data import Dataset, DataLoader
import torch.nn.functional as F

%matplotlib inline
%config InlineBackend.figure_format='retina'

sns.set(style='whitegrid', palette='muted', font_scale=1.2)

HAPPY_COLORS_PALETTE = ["#01BEFE", "#FFDD00", "#FF7D00", "#FF006D", "#ADFF02", "#8F00FF"]

sns.set_palette(sns.color_palette(HAPPY_COLORS_PALETTE))

rcParams['figure.figsize'] = 12, 8

RANDOM_SEED = 42
np.random.seed(RANDOM_SEED)
torch.manual_seed(RANDOM_SEED)

device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
device

device(type='cuda', index=0)
```

```
####!gdown --id 1S6qMioqPJjyBLpLVz4gmRTnJHnjitnuV
####!gdown --id 1zdmewp7ayS4js4VtrJEHzAheSW-5NBZv
```

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
print(testing_data.columns, training_data.columns)
```

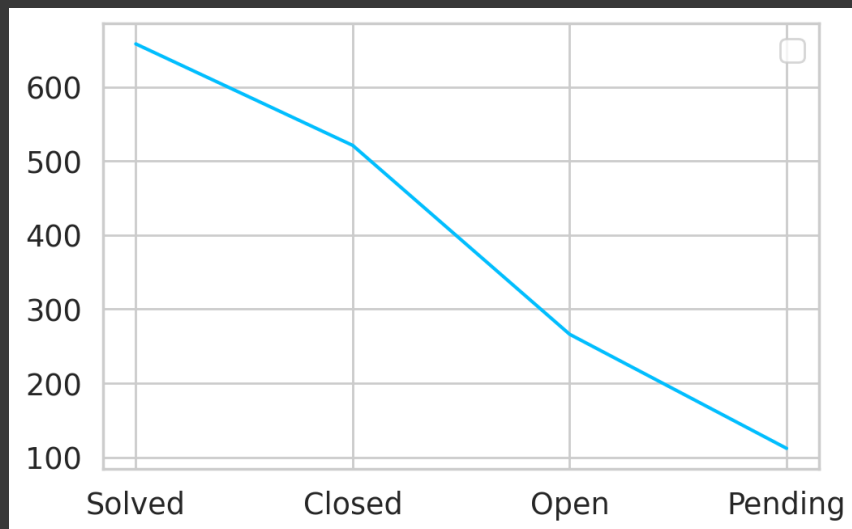
```
Index(['Ticket #', 'Customer Complaint', 'Date', 'Date_month_year', 'Time',
      'Received Via', 'City', 'State', 'Zip code', 'Status',
      'Filing on Behalf of Someone'],
      dtype='object') Index(['Ticket #', 'Customer Complaint', 'Date', 'Date_month_year', 'Time',
      'Received Via', 'City', 'State', 'Zip code', 'Status',
      'Filing on Behalf of Someone'],
      dtype='object')
```

```
training_data.Status.value_counts()
```

```
Solved      658
Closed      521
Open        266
Pending     112
Name: Status, dtype: int64
```

```
plt.rcParams["figure.figsize"] = [5.50, 3.50]
plt.rcParams["figure.autolayout"] = True
```

```
training= training_data.Status.value_counts()
labels = ['Solved', 'Closed', 'Open', 'Pending']
plt.legend(labels, loc="best")
plt.plot(training)
plt.show()
```



```
print(training_data.columns)
```

```
Index(['Ticket #', 'Customer Complaint', 'Date', 'Date_month_year', 'Time',
      'Received Via', 'City', 'State', 'Zip code', 'Status',
      'Filing on Behalf of Someone'],
      dtype='object')
```

```
training_data["length"]=training_data['Customer Complaint'].apply(len)
```

```
testing_data["length"]=testing_data['Customer Complaint'].apply(len)
```

```
import string
import re
def clean_text(text):
    '''Make text lowercase, remove text in square brackets,remove links,remove punctuation
    and remove words containing numbers.'''
    text = text.lower()
    text = re.sub('\[.*?\]', '', text)
    text = re.sub('https?://\S+|www\.\S+', '', text)
    text = re.sub('<.*?>+', '', text)
```

```
text = re.sub('[%s]' % re.escape(string.punctuation), '', text)
text = re.sub('\n', '', text)
text = re.sub('\w*\d\w*', '', text)
return text
```

```
training_data['Customer Complaint'] = training_data['Customer Complaint'].apply(lambda x: clean_text(x))
```

```
testing_data['Customer Complaint'] = testing_data['Customer Complaint'].apply(lambda x: clean_text(x))
```

```
import nltk
nltk.download('stopwords')
from nltk.corpus import stopwords
stop = stopwords.words('english')
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

```
training_data['Clean Customer Complaint'] = training_data['Customer Complaint'].apply(lambda x: " ".join(x for x in x.split() if x))
training_data['Clean Customer Complaint'].head()
```

```
1236                                comcast uverse
541    billingcontract issue related data cap
539                                inaccurate billing
1566    comcast fraudulent billing charges
632    comcast east windsor nj complaint
Name: Clean Customer Complaint, dtype: object
```

```
testing_data['Clean Customer Complaint'] = testing_data['Customer Complaint'].apply(lambda x: " ".join(x for x in x.split() if x))
testing_data['Clean Customer Complaint'].head()
```

```
2                                speed service
4    comcast working service boot
5    isp charging arbitrary data limits overage fees
10   billing service asked disconnected
11   yahoo failure restore email search feature
Name: Clean Customer Complaint, dtype: object
```

```
###!pip install nltk
```

```
import nltk
```

```
nltk.download('punkt')
```

```
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Package punkt is already up-to-date!
True
```

```
from nltk.stem.snowball import SnowballStemmer
```

```
stemmer = nltk.stem.SnowballStemmer('english')
```

```
nltk.download('stopwords')
stop_words = set(nltk.corpus.stopwords.words('english'))
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

```
import re
```

```
def preprocessing(text):
    tokens = [word for word in nltk.word_tokenize(text) if (len(word) > 3 and len(word.strip('Xx/')) > 2 and len(re.sub('\d+', '')) > 2)]
    tokens = map(str.lower, tokens)
    stems = [stemmer.stem(item) for item in tokens if (item not in stop_words)]
    return stems
```

```
from sklearn.feature_extraction.text import TfidfVectorizer
```

```
vectorizer_tf = TfidfVectorizer(tokenizer=preprocessing, stop_words=None, max_df=0.75, max_features=1000, lowercase=False, ngram_range=(1, 2))
```

```
print(testing_data.columns, training_data.columns)
```

```
Index(['Ticket #', 'Customer Complaint', 'Date', 'Date_month_year', 'Time',
       'Received Via', 'City', 'State', 'Zip code', 'Status',
       'Filing on Behalf of Someone', 'length', 'Clean Customer Complaint'],
      dtype='object') Index(['Ticket #', 'Customer Complaint', 'Date', 'Date_month_year', 'Time',
       'Received Via', 'City', 'State', 'Zip code', 'Status',
       'Filing on Behalf of Someone', 'length', 'Clean Customer Complaint'],
      dtype='object')
```

```
training_data['encoded_Status'] = training_data['Status'].astype('category').cat.codes
testing_data['encoded_Status'] = testing_data['Status'].astype('category').cat.codes
```

```
training_data['encoded_Status'].value_counts()
```

```
3      658
0      521
1      266
2      112
Name: encoded_Status, dtype: int64
```

```
####! pip install transformers
```

```
from transformers import DistilBertTokenizerFast
from transformers import TFDistilBertForSequenceClassification
import tensorflow as tf
import pandas as pd
import json
import gc
```

```
traindata_texts = training_data["Clean Customer Complaint"].to_list() # Features (not-tokenized yet)
traindata_labels = training_data["encoded_Status"].to_list() # Labels
```

```
testdata_texts = testing_data["Clean Customer Complaint"].to_list() # Features (not-tokenized yet)
testdata_labels = testing_data["encoded_Status"].to_list() # Labels
```

```
tokenizer = DistilBertTokenizerFast.from_pretrained('distilbert-base-uncased')
train_encodings = tokenizer.batch_encode_plus(traindata_texts, truncation=True, padding=True)
test_encodings = tokenizer.batch_encode_plus(testdata_texts, truncation=True, padding=True)
```

```
traindata_labels = tf.keras.utils.to_categorical(traindata_labels, dtype="uint8")
testdata_labels = tf.keras.utils.to_categorical(testdata_labels, dtype="uint8")
```

```
testdata_texts
```

```

    'unfair billing practices',
    'comcast complaint',
    'comcast bars hbogo streaming purchased device choosing',
    'complaint comcast',
    'comcast inaccurately measuring bandwidth consumption',
    'comcast unlimited internet access recently limited per month',
    'data cap',
    'comcast service billing',
    'increased monthly fee year contract',
    'shitty comcast service',
    'terrible pricing one viable option',
    'pay cable line',
    'access service',
    'comcast xfinity',
    'comcastxfinity communication',
    'comcastnon existent customer service terrible internet connection',
    'internet service loses signal',
    'comcast bill',
    'inexplicable disconnection subsequent mishandling',
    'false sales offers',
    'paying getting',
    'throttling comcast internet services',
    'comcast excessive charge cancelling service',
    'comcast refusal service',
    'deceptive trade practices false advertising bait switch',
    'sold billed short term services',
    'comcast transfer service complaint',
    'unauthorized charges account comcast',
    'comcast charge router unreturned equipment',
    'egregious fees',
    'cramming false internet promotion speed complaint',
    'comcast business internet',
    'comcast',
    'comcast xfinity',
    'comcast sales people reflect said bill',
    'slow internet',
    'comcast cramming',
    'billing continues months terminating service',
    'data usage caps',
    'comcast charges',
    'monopolistic billing practices',
    'comcast outage',
    'fed comcast',
    'comcast charge',
    'comcast xfinity internet data caps',
    'comcast data cap',
    'inconsistent intermittent internet connectivity',
    'comcast internet',
    'comcast cap',
    'false advertisingbait switch',
    'comcast customer service billing issues',
    'sent check payment comcast',
    'comcast unfair pricing',
    'comcast monthly billing returned modem',
    'extremely unsatisfied comcast customer']

```

```

train_dataset = tf.data.Dataset.from_tensor_slices((
    dict(train_encodings),
    traindata_labels
))
test_dataset = tf.data.Dataset.from_tensor_slices((
    dict(test_encodings),
    testdata_labels
))

```

```

#metrics = [tf.keras.metrics.SparseCategoricalAccuracy('accuracy', dtype=tf.float32)]
metrics = tf.metrics.CategoricalAccuracy()

```

```

loss = tf.keras.losses.CategoricalCrossentropy(from_logits=True)

```



```
model = TFDistilBertForSequenceClassification.from_pretrained('distilbert-base-uncased', num_labels=4)
```

```
optimizer = tf.keras.optimizers.Adam(learning_rate=5e-5)  
model.compile(optimizer=optimizer, loss=loss, metrics=['accuracy'])
```

Some weights of the PyTorch model were not used when initializing the TF 2.0 model TFDistilBertForSequenceClassification:  
- This IS expected if you are initializing TFDistilBertForSequenceClassification from a PyTorch model trained on another task.  
- This IS NOT expected if you are initializing TFDistilBertForSequenceClassification from a PyTorch model that you expect to be initialized from the TF 2.0 model TFDistilBertForSequenceClassification.  
Some weights or buffers of the TF 2.0 model TFDistilBertForSequenceClassification were not initialized from the PyTorch model.  
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
model.fit(train_dataset.shuffle(800).batch(24), epochs=30, batch_size=24,  
          validation_data=test_dataset.shuffle(800).batch(24))
```

```
Epoch 3/30  
65/65 [=====] - 7s 109ms/step - loss: 1.1151 - accuracy: 0.4611 - val_loss: 1.1690 - val_accuracy: 0.4611  
Epoch 4/30  
65/65 [=====] - 6s 90ms/step - loss: 1.0133 - accuracy: 0.5427 - val_loss: 1.1847 - val_accuracy: 0.5427  
Epoch 5/30  
65/65 [=====] - 6s 95ms/step - loss: 0.8730 - accuracy: 0.6191 - val_loss: 1.3289 - val_accuracy: 0.6191  
Epoch 6/30  
65/65 [=====] - 5s 80ms/step - loss: 0.7103 - accuracy: 0.7103 - val_loss: 1.5809 - val_accuracy: 0.7103  
Epoch 7/30  
65/65 [=====] - 6s 91ms/step - loss: 0.6274 - accuracy: 0.7360 - val_loss: 1.6180 - val_accuracy: 0.7360  
Epoch 8/30  
65/65 [=====] - 8s 127ms/step - loss: 0.5601 - accuracy: 0.7662 - val_loss: 1.6508 - val_accuracy: 0.7662  
Epoch 9/30  
65/65 [=====] - 7s 102ms/step - loss: 0.4852 - accuracy: 0.8125 - val_loss: 1.8595 - val_accuracy: 0.8125  
Epoch 10/30  
65/65 [=====] - 6s 84ms/step - loss: 0.4421 - accuracy: 0.8157 - val_loss: 2.0077 - val_accuracy: 0.8157  
Epoch 11/30  
65/65 [=====] - 6s 96ms/step - loss: 0.4161 - accuracy: 0.8170 - val_loss: 2.0964 - val_accuracy: 0.8170  
Epoch 12/30  
65/65 [=====] - 6s 97ms/step - loss: 0.3916 - accuracy: 0.8343 - val_loss: 2.1156 - val_accuracy: 0.8343  
Epoch 13/30  
65/65 [=====] - 6s 91ms/step - loss: 0.3709 - accuracy: 0.8343 - val_loss: 2.2491 - val_accuracy: 0.8343  
Epoch 14/30  
65/65 [=====] - 5s 85ms/step - loss: 0.3591 - accuracy: 0.8330 - val_loss: 2.2467 - val_accuracy: 0.8330  
Epoch 15/30  
65/65 [=====] - 6s 89ms/step - loss: 0.3335 - accuracy: 0.8478 - val_loss: 2.3809 - val_accuracy: 0.8478  
Epoch 16/30  
65/65 [=====] - 8s 123ms/step - loss: 0.3277 - accuracy: 0.8497 - val_loss: 2.5650 - val_accuracy: 0.8497  
Epoch 17/30  
65/65 [=====] - 5s 82ms/step - loss: 0.3328 - accuracy: 0.8388 - val_loss: 2.3242 - val_accuracy: 0.8388  
Epoch 18/30  
65/65 [=====] - 6s 95ms/step - loss: 0.3448 - accuracy: 0.8465 - val_loss: 2.6136 - val_accuracy: 0.8465  
Epoch 19/30  
65/65 [=====] - 5s 84ms/step - loss: 0.3193 - accuracy: 0.8516 - val_loss: 2.6151 - val_accuracy: 0.8516  
Epoch 20/30  
65/65 [=====] - 6s 95ms/step - loss: 0.3087 - accuracy: 0.8536 - val_loss: 2.7352 - val_accuracy: 0.8536  
Epoch 21/30  
65/65 [=====] - 5s 81ms/step - loss: 0.3128 - accuracy: 0.8478 - val_loss: 2.7212 - val_accuracy: 0.8478  
Epoch 22/30  
65/65 [=====] - 5s 84ms/step - loss: 0.3034 - accuracy: 0.8600 - val_loss: 2.7213 - val_accuracy: 0.8600  
Epoch 23/30  
65/65 [=====] - 5s 84ms/step - loss: 0.3107 - accuracy: 0.8504 - val_loss: 2.6868 - val_accuracy: 0.8504  
Epoch 24/30  
65/65 [=====] - 6s 88ms/step - loss: 0.2979 - accuracy: 0.8516 - val_loss: 2.8850 - val_accuracy: 0.8516  
Epoch 25/30  
65/65 [=====] - 6s 86ms/step - loss: 0.3185 - accuracy: 0.8548 - val_loss: 2.7074 - val_accuracy: 0.8548  
Epoch 26/30  
65/65 [=====] - 5s 81ms/step - loss: 0.3073 - accuracy: 0.8459 - val_loss: 2.7514 - val_accuracy: 0.8459  
Epoch 27/30  
65/65 [=====] - 6s 95ms/step - loss: 0.2944 - accuracy: 0.8542 - val_loss: 2.8170 - val_accuracy: 0.8542
```

```
Epoch 30/30
65/65 [=====] - 5s 81ms/step - loss: 0.2890 - accuracy: 0.8529 - val_loss: 2.6823 - val_accuracy: 0.4273
<keras.callbacks.History at 0x7aad58307cd0>
```

```
# the instantiated 🤗 Transformers model - Distilled Bert
```

```
from transformers import TFDistilBertForSequenceClassification, TFTTrainer, TFTTrainingArguments
```

```
model.evaluate(test_dataset.shuffle(20).batch(12),
               return_dict=True,
               batch_size=12)
```

```
56/56 [=====] - 1s 19ms/step - loss: 2.6823 - accuracy: 0.4273
{'loss': 2.6823196411132812, 'accuracy': 0.4272863566875458}
```

```
y_pred = model.predict(test_dataset)
```

```
667/667 [=====] - 5s 6ms/step
```

```
predictions = tf.nn.softmax(y_pred.logits)
# Extracting the indices with the highest probabilities
predictions = tf.argmax(predictions, axis=1).numpy()
```

```
predictions = pd.DataFrame(predictions)
```

```
predictions.value_counts()
```

```
3    347
0    211
1     94
2     15
dtype: int64
```

```
predictions.rename(columns = {0:"Predict"}, inplace=True)
```

```
train2 = pd.read_csv("/content/train.csv")
```

```
test2 = pd.read_csv("/content/test.csv")
```

```
test2["Status"].value_counts()
```

```
Solved    315
Closed    213
Open       97
Pending    42
Name: Status, dtype: int64
```

```
test2["Status"].value_counts()
```

```
Solved    315
Closed    213
Open       97
Pending    42
Name: Status, dtype: int64
```

```
class_names = {
    3 : "Solved",
    0 : "Closed",
    1 : "Open",
    2 : "Pending"
```

```
}
```

```
predictions.value_counts()
```

```
Predict
3      347
0      211
1       94
2       15
dtype: int64
```

```
predictions = predictions['Predict'].map(class_names)
```

```
predictions.value_counts()
```

```
Solved      347
Closed      211
Open         94
Pending      15
Name: Predict, dtype: int64
```

✓ 0s completed at 11:20 PM

